



Menu

Pricing

Simple and flexible. Only pay for what you use.

[Contact sales](#)[Learn more ↓](#)



Language models

Multiple models, each with different capabilities and price points. Prices are per 1,000 tokens. You can think of tokens as pieces of words, where 1,000 tokens is about 750 words. This paragraph is 35 tokens.



natural language and solve difficult problems with accuracy.

[Learn about GPT-4](#)

8K context

+

32K context

+

GPT-3.5 Turbo

GPT-3.5 Turbo models are capable and cost-effective.

`gpt-3.5-turbo` is the flagship model of this family and is optimized for dialog.

`gpt-3.5-turbo-instruct` is an Instruct model and only supports a 4K context window.

[Learn about GPT-3.5 Turbo ↗](#)

4K context

+

16K context

+

Fine-tuning models

Create your own custom models by fine-tuning our base models with your training data. Once you fine-tune a model, you'll be billed only for the tokens you use in requests to that model.

[Learn about fine-tuning ↗](#)

babbage-002

+

davinci-002

+

GPT-3.5 Turbo

+



embeddings offering.

[Learn about embeddings ↗](#)

Ada v2

+

Base models

GPT base models are not optimized for instruction-following and are less capable, but they can be effective when fine-tuned for narrow tasks.

[Learn about GPT base models ↗](#)

babbage-002

+

davinci-002

+

Other models

Image models

Build DALL-E directly into your apps to generate and edit novel images and art. Our image models offer three tiers of resolution for flexibility.

[Learn about image generation ↗](#)

1024×1024

+

512×512

+

256×256

+



[Learn about Whisper ↗](#)

Whisper

+

Older models

We continue to improve our models and periodically retire older, less used models.

[View pricing and info for older models ↗](#)

Usage quotas

When you sign up, you'll be granted an initial spend limit, or quota, and we'll increase that limit over time as you build a track record with your application. If you need more tokens, you can always request a quota increase.

[Request quota increase ↗](#)

Simple and flexible

Start for free

Start experimenting with \$5 in free credit that can be used during your first 3 months.

Pay as you go

To keep things simple and flexible, pay only for the resources you use.

Choose your model

Use the right model for the job. We offer a spectrum of capabilities and price points.



Morgan Stanley

Morgan Stanley wealth management deploys GPT-4 to organize its vast knowledge base.



stripe DOCS



Ask a question or search the docs...

Stripe Docs

Explore our guides and examples to integrate Stripe.

[Get started with payments >](#)[Explore all products](#)

RECENTLY VIEWED

[Accepting online payments](#)
[Subscriptions](#)

API KEYS

Publishable key pk_abc_123_***
Secret key sk_xyz_123_***

Test cards →

4242 4242 4242 4242

No-code

[Sell and get paid online](#)
[Invoice clients and customers](#)
[Set up recurring payments](#)

Stripe hosted

[Create a hosted invoice page](#)
[Set up the customer portal](#)
[Create a pricing table](#)

For developers

[API reference](#)
[Stripe.js reference](#)
[Set up the CLI](#)

Test the Stripe API

[Start a payment](#)[Open an invoice](#)[Onboard platform account](#)[Generate a payment link](#) No code[Create a subscription](#) No code

```
$ stripe payment_intents create --amount 1099 --currency "usd"
```

Browse by product

DIRECT PAYMENTS

Payments
Online payments

FINANCIAL SERVICES

Capital
Business financing

BUSINESS OPERATIONS

Radar
Fraud & risk management

Stripe

Stripe leverages GPT-4 to streamline user experience and combat fraud.



Wix

Deloitte.



What's a token? +

Which model should I use? +

How will I know how many tokens I've used each month? +

How can I manage my spending? +

Is the ChatGPT API included in the ChatGPT Plus subscription? +

Does Playground usage count against my quota? +

How is pricing calculated for Completions? +

How is pricing calculated for Fine-tuning? +

Is there an SLA on the various models? +



Start creating with OpenAI's powerful models.

[Get started ↗](#)

[Contact sales](#)



Research

[Overview](#)

[Index](#)

[GPT-4](#)

[DALL-E 3](#)

API

[Overview](#)

[Data privacy](#)

[Pricing](#)

[Docs ↗](#)

ChatGPT

[Overview](#)

Company

[About](#)

[About](#)[Security](#)[Customer stories](#)[Safety](#)

OpenAI © 2015–2023[Terms & policies](#)[Privacy policy](#)[Brand guidelines](#)**Social**[Twitter](#)[YouTube](#)[GitHub](#)[SoundCloud](#)[LinkedIn](#)[Back to top ↑](#)

